Music genre classification

Emanuele Pocelli, Fabrizio Sordetti

June 26, 2024

Contents

1 Introduction

This is the report of the NAML project 2023/2024, which consists of reproducing the results of a paper on classification of music genres.

2 The problem

The problem consists of classifying songs according to their genre. The given paper proposes a novel architecture to solve this problem efficiently and with high performance. It is based on convolutional neural networks. As is often the case when processing audios using neural networks, we work on the spectrogram of the audio and treat It as an image. This is because architectures for images are more advanced, since the interest in image processing is bigger.

More in detail, the architecture is composed of two pieces: a unet followed by parallel convolutions. The unet is a recent architecture that features an encoder and a decoder: we get fine grained features in the contracting path, where we downsample while increasing the number of filters in each step, then, in the expansive path, we upsample while decreasing the number of filters in each step, also using the skip features from the contracting path. Such architecture excels in image segmentation. Finally, parallel convolutions are used to analyse the fine grained features.

In the paper, multiple music genre classification datasets are used, but we will focus on the GTZAN that is the most famous and probably the most interesting.


Figure 1: Plot of the unet model. We can see the U-shape, hence the name

Figure 2: Parallel convolutions

3   The dataset

The dataset used is the GTZAN, which is widely available online. It consists in 1000 songs belonging to 10 different genres. For each song, there is a 30s audio clip. This dataset is often used to test new architectures. It is quite complex, so strong models are needed in order to achieve a decent performance.
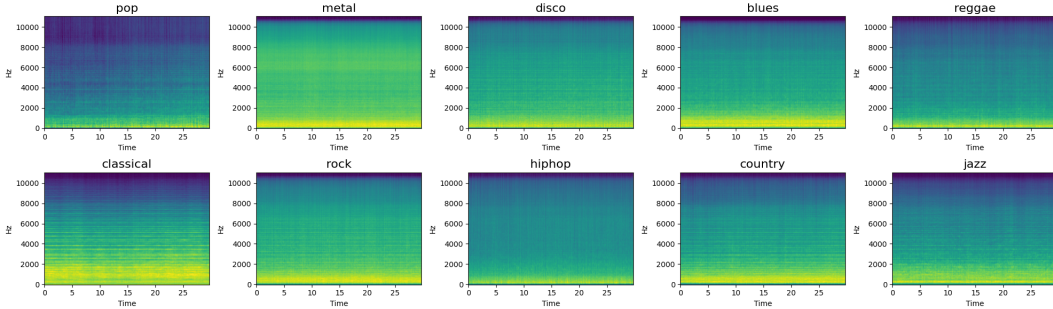

Figure 3: Mean spectrograms for each genre. We can see clear differences for some genres, however for others are much less noticeable

4   How we approached the problem

The architecture proposed in the paper has a size of more than 300MB. We had significant limitations on the computational side, both in time and space, so we could not use the very same architecture.

Our architecture is based on the one proposed in the paper, but It is revised and much smaller (15MBs). We tried training bigger models, but a lot of problems arised, like strong overfitting or exploding gradients. Very special care was needed in order to avoid these phenomena, and, as we will see in the next sections, we were hindered by our computational limits and could not overcome these problems because of that.

A significant modification that was necessary has been the use of ReLU function instead of PReLU. While It is true that PReLU gives the model a higher degree of freedom and prevents the problem of vanishing gradient, It also makes the training much less stable and often causes the opposite problem, exploding gradient.

We used batch normalization, l2 regularization and dropout. The paper did not mention these aspects, but we had a much more stable training with them. It has been a bit tricky to find out "where" to place each of these layers, because CNNs are very sensitive in this regard.


Figure 4: Plot of the model using visualkeras

5   Data processing

When dealing with deep neural networks, especially when trainin them from scratch, we need a lot of data. Unfortunately, the GTZAN dataset is relatively small, only featuring 1000 songs divided into 10 genres. Because of that, the data processing was critical in order to avoid overfitting.
Even though the dataset is small for a NN, It is pretty big for the RAM: we faced a lot of crashes because of memory and we were forced to stream the dataset during training using TensorFlow's data generator class. This, however, made the process slower.

First of all, we split each song into 5 different pieces so that each input is 6s long. This is an easy way to have a bigger dataset while not losing too much information: It is still possible to classify the music genre given only a 6s clip. Every clip of the same song all belong either to the training set or the validation set. This is very important to keep the independence between the two datasets, otherwise It would be like cheating.
All the data was normalised to be in the range of (0,1).

The key part is data augmentation. The paper talks very little about It: they simply state that they implement It using standard librosa functions.
Augmentation on a spectrogram is a bit tricky and It is not possible to use general techniques designed for images, like rotation or mirroring.
There are techniques that work on the waveform or on the spectrogram Itself. The former provides a stronger effect, but is computationally very expensive, since It requires to generate the spectrogram of each song in every epoch. Librosa functions are available for this purpose, which is why we believe this was the choice of the authors.
The latter is weaker, but since It works on the spectrogram Itself, It is quite cheap to perform.
We implemented an augmentation pipeline using both approaches, but because of computational constraints, we opted for the second option in the end. This made the training feasible for us at the cost of more overfitting.
We used the following techniques: random noise, time shift, time mask, frequency mask.
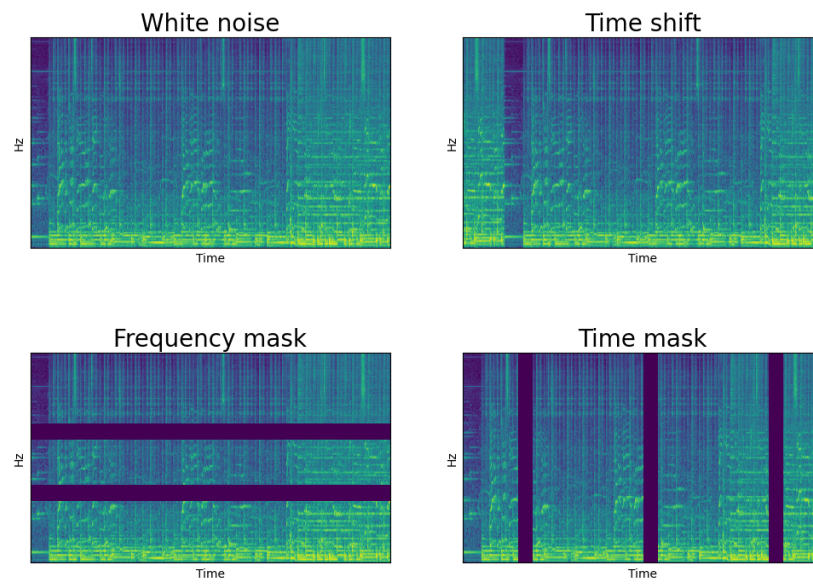

Figure 5: Example of a spectrogram

Figure 6: The different techniques we used

Figure 7: All the techniques combined

6 Results

We were pretty happy with the results we got. There is a bit of overfitting, but overall they are decent considering the size of our model and the complexity of the problem.
The final accuracy achieved was 76% in the training set and 66.7% in the validation set.

Figure 8: Accuracy and loss in each epoch

Figure 9: Confusion matrix

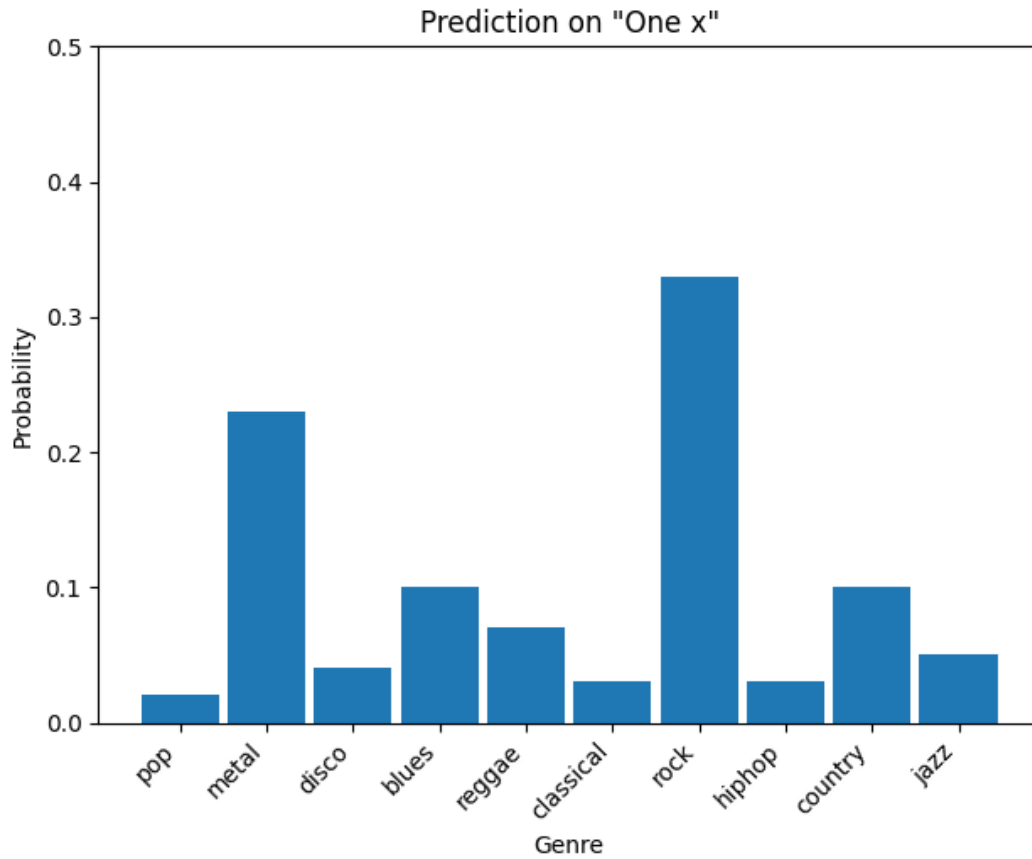Figure 10: Statistical results for each genre

Figure 11: Example output on a song

The images contain the aggregated results (training+validation). In the images folder, there are the images for training validation separately. The performance is slightly lower on validation and slightly higher on training because of overfitting.

It is relevant to see that the confusion matrix is overall pretty good. Most of the error came from the misclassification of reggae and country with metal. Looking back at the mean spectrograms, country, reggae and metal are very similar.

Instead, we were surprised to not see many misclasifications between rock and metal, since the genres sound and look in the spectrograms very similar.

On the other hand, jazz and classical music have the most unique spectrograms when compared to the other genres. This explains why they have the highest accuracy, with low false positives and false negatives.

7   Conclusions

The paper is not very clear about their results. Focusing on accuracy, at first the authors report 99.4%, then, in the next figure, they state a value that ranges from 99% to 100%, with an average of 99.85%, however in the reported confusion matrix, the accuracy appears to be 97%.
We are not sure why the results are so chaotic. Also, an accuracy as high as 99.85% in this complex problem may be a symptom of overfitting, and It may be possible that the independence between training set and validation set was not preserved (for instance if a song was split into 5 clips, with some clips going to the training set and others in the validation set).
In any case, we did not get really close to their performance. It is still remarkable for a model so small to achieve such a high performance in this complex task.
Computational limitations unfortunately hindered our efforts. Having access to more resources to speed up the training would allow us to implement a stronger augmentation to overcome overfitting and deploy a bigger model that would maybe achieve better performance.
To sum up, the architecture is definetely capable and promising, even though we had to interpret a good part of It. the paper was not very clear in many aspects. Difficulties in the language are a possible reason for It.